

PCSE Spring 2016 Homework 4: Exchange-Sort Exercise

Neil Klenk - nlk322

2/23/16

0.1 Introduction

In this exercise, the exchange-sort algorithm was implemented between many processors. Each processor was given a single number, and compared with its neighboring processor to determine if their values needed to be swapped or not. The algorithm was split into two stages: Even stage: processors $2i$ and $2i + 1$ compare and swap. Odd stage: processors $2i + 1$ and $2i + 2$ compare and swap. by repeating this $P/2$ times, where P is the number of processors, the processor numbers in ascending order would contain the corresponding number in ascending order.

0.2 Theory

1. In total the parallel code should take $P/2$ iterations where P is the number of processors being used. 2. The algorithm takes $P/2$ sequential steps. T_1 : 0.00549 sec. This

each comparison on the initial node took approximately 0.000011 seconds. For an array consisting of 32 values, it would take 0.00549 seconds to perform all of those comparisons on one processor. T_p : 0.067 sec for 32 processors (2 nodes). $T(\infty)$:

0.000177 sec. Since each processor only contains one value, and this comparison is performed directly between a processor and its partner, there is no opportunity to gain parallelism once you leave the first node. S_p : 0.082 sec. This was determined by dividing T_1 by the actual time it took (0.067 sec). This low value demonstrated the effects of network latency. E_p : 0.002563 sec. Average amount of parallelism: 31.0169. This value was obtained by dividing the time it would take on one processor by the amount of time it took when constrained to one node. 3. T_1 is the sequential time (Time for

one processor to work though it). Thus S_p and E_p would remain the same.

0.3 Writeup

0.3.1 The Algorithm and Corresponding Code

This section of the code takes care of the edge cases that appear when the partner rank is set to either -1 or "size". This occurs in the odd stage where the partner to rank = 0 is rank = -1, when the partner to rank = 31 is rank = 32. In each case the processor does not exist, so they send and receive from and to "MPI.PROC_NULL" to prevent the code from hanging.

```
if(rank == 0)
prevProc = MPI_PROC_NULL;
else
```

```

prevProc = rank - 1;

if(rank == size - 1)
nextProc = MPI_PROC_NULL;
else
nextProc = rank + 1;

```

This section of the code is the meat of the program. It functions by first evaluating the stage that the program is in, either even or odd. Depending on if it is the even or odd stage, the swap and subsequent compare happen between different processors. If it is the even stage, the even processor compares with the following processor. The odd processor keeps the larger of the two values while the even processor keeps the smaller of the two values. If it is in the odd stage, the even processor compares with the preceding processor. The even processor keeps the larger value while the odd processor keeps the smaller value.

```

while(stage_count <= size/2){
other_num = process_num;
  if(stage_count % 2 == 0){ //even stage
    if(rank % 2 == 0){ //if even
      MPI_Send(&process_num,1,MPI_DOUBLE,nextProc,10,MPI_COMM_WORLD);
      MPI_Recv(&other_num,1,MPI_DOUBLE,nextProc, 20,
        MPI_COMM_WORLD, MPI_STATUS_IGNORE);
      process_num = (process_num < other_num) ? process_num:
        other_num;
    }
    else{ //if odd
      MPI_Recv(&other_num,1,MPI_DOUBLE,prevProc, 10,
        MPI_COMM_WORLD, MPI_STATUS_IGNORE);
      MPI_Send(&process_num,1,MPI_DOUBLE,prevProc,20,MPI_COMM_WORLD);
      process_num = (process_num > other_num) ? process_num:
        other_num;
    }
  }
}
else{ //odd stage
  if(rank % 2 != 0){ //if odd
    MPI_Send(&process_num,1,MPI_DOUBLE,nextProc,30,MPI_COMM_WORLD);
    MPI_Recv(&other_num,1,MPI_DOUBLE,nextProc, 40,
      MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    process_num = (process_num < other_num) ? process_num:
      other_num;
  }
  else{ //if even
    MPI_Recv(&other_num,1,MPI_DOUBLE,prevProc, 30,
      MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    MPI_Send(&process_num,1,MPI_DOUBLE,prevProc,40,MPI_COMM_WORLD);
  }
}

```

```
        process_num = (process_num > other_num) ? process_num :  
            other_num;  
    }  
}  
stage_count++;  
}
```

0.3.2 Theoretical Issues

Theoretical issues with this implementation include the potential for the MPI_Sendrecv be effected by its blocking properties.

0.3.3 Timing Results

```
Rank: 15 Time: 0.067246
Rank: 11 Time: 0.067265
Rank: 3 Time: 0.067229
Rank: 13 Time: 0.067323
Rank: 7 Time: 0.067190
Rank: 10 Time: 0.067301
Rank: 9 Time: 0.067232
Rank: 5 Time: 0.067189
Rank: 1 Time: 0.067185
Rank: 2 Time: 0.067229
Rank: 4 Time: 0.067196
Rank: 6 Time: 0.067200
Rank: 14 Time: 0.067247
Rank: 12 Time: 0.067257
Rank: 31 Time: 0.000090
Rank: 30 Time: 0.000100
Rank: 29 Time: 0.000168
Rank: 27 Time: 0.000136
Rank: 26 Time: 0.000134
Rank: 25 Time: 0.000126
Rank: 28 Time: 0.000128
Rank: 24 Time: 0.000125
Rank: 17 Time: 0.000158
Rank: 23 Time: 0.000136
Rank: 22 Time: 0.000128
Rank: 18 Time: 0.000177
Rank: 20 Time: 0.000134
Rank: 19 Time: 0.000175
Rank: 16 Time: 0.000153
Rank: 21 Time: 0.000147
Rank: 8 Time: 0.067240
Rank: 0 Time: 0.067230
```

The two concentrations of times indicates that some of the processors had to communicate across the nodes. Inter node communications are much faster than intra node communication, and explains why there is a concentration of times at both 0.00017 and 0.067.